

Fused Block-Sparse Attention with Learned Content-Dependent Selection:

A Triton Kernel for Memory-Efficient Long-Context Inference

Miroslav Drbal
mdrbal@nymfe.net

Gen Digital, Inc. | miroslav.drbal@gendigital.com

Abstract

We present a block-sparse attention mechanism with a fused Triton kernel that achieves memory parity with dense FlashAttention while delivering up to $7.4\times$ latency speedup at 262k-token context length. The block selector is trained jointly with the model via MoE-style gate coupling and an auxiliary routing loss, reaching near-perfect block-retrieval accuracy (hit = 0.99) without a separate pre-training stage. On a 45M-parameter language model trained on TinyStories, the sparse model achieves validation perplexity 15.0 versus 14.4 for the dense baseline, a 4% gap that narrows to parity at shorter context windows. We also investigate centroid-based linear selection ($\mathcal{O}(N \cdot C)$) as a replacement for the $\mathcal{O}(N^2/B_s^2)$ block-pair selection, finding that InfoNCE contrastive training fixes a critical collapse failure in the routing loss (hit $0.19 \rightarrow 1.0$) but that task accuracy remains bottlenecked by selection precision, not routing recall. We further isolate selection from the attention read up to 16M tokens and show that single-level block-pair selection has an $\mathcal{O}(N^{4/3})$ compute floor: no block-size schedule yields end-to-end linearity, as the quadratic only moves between the selection and read stages. A second, lossless linear selector, LSH bucketing, matches block-sparse accuracy under supervised routing (0.92 vs. 0.96 on MQAR) and far exceeds centroid, but its unsupervised (gate-only) routing is inert in a 45M language model, matching a no-routing baseline, leaving reliable unsupervised linear selection open. All experiments run on a single NVIDIA GB10 Grace-Blackwell GPU.

1 Introduction

Transformer-based language models face a quadratic attention bottleneck: standard softmax attention scales as $\mathcal{O}(N^2)$ in sequence length N , limiting practical context windows. While FlashAttention [1, 2] reduces the *constant* factor by tiling computations in SRAM, the asymptotic cost remains quadratic.

Content-dependent sparse attention offers a path to genuine sub-quadratic scaling by selecting a small subset of key/value blocks per query block. Mechanisms such as Native Sparse Attention (NSA) [3], MoBA [4], and the Routing Transformer [5] learn to route queries to relevant key blocks, reducing the attention cost to $\mathcal{O}(N \cdot k \cdot B_s)$ where k is the number of selected blocks and B_s is the block size.

However, two practical barriers remain:

1. **Selector training.** Discrete top- k selection severs the gradient path to the selector parameters. Without auxiliary supervision, the selector remains a frozen random projection.
2. **Gather overhead.** Even with correct selection, the reference PyTorch implementation materializes gathered key/value tensors, consuming up to $40\times$ more memory than dense FlashAttention and negating the theoretical savings.

We address both barriers. For (1), we use MoE-style gate coupling [6] -adding log-sigmoid(score) as a bias to attention logits, so the selector receives task gradients. An auxiliary routing loss (cross-entropy on the known needle block for synthetic tasks, InfoNCE contrastive loss [14] for centroid routing) provides cold-start supervision. For (2), we implement a fused Triton kernel [13] that reads selected block indices, gathers key/value blocks one at a time from the KV cache, and computes online-softmax in SRAM, never materializing the gathered tensor.

Contributions.

- A **fused Triton kernel** for block-sparse attention (forward + backward) achieving memory parity with dense FlashAttention (1.11 GB vs. 1.12 GB at 262k tokens) and $7.4\times$ latency speedup (Section 4).
- **Gate-coupled selector training** with InfoNCE routing loss, fixing a collapse failure in centroid-based linear selection (hit $0.19 \rightarrow 1.0$) and characterizing the precision/recall tradeoff (Section 5).
- An **end-to-end language model** (45M params, TinyStories [12]) trained with the fused kernel, achieving val perplexity 15.0 vs. 14.4 dense (Section 6).
- A **negative result**: linear-selection centroid routing achieves perfect retrieval recall but cannot match block-sparse task accuracy, identifying representation quality under coarse routing as the bottleneck (Section 5).
- A **scaling analysis** isolating selection from the attention read to 16M tokens: the read is linear, but single-level block-pair selection has an $\mathcal{O}(N^{4/3})$ floor that no block-size schedule can beat and that exhausts memory before any latency crossover (Section 7.1).

2 Background and Related Work

Sparse attention. Fixed-pattern sparsity (Longformer [7], BigBird [8]) restricts attention to local windows plus global tokens but cannot solve content-dependent retrieval tasks like multi-query associative recall (MQAR) [15, 11]. Content-dependent methods -Routing Transformer [5], Reformer [9], Sparse Sinkhorn Attention [10], NSA [3], MoBA [4] -learn to select relevant blocks but face the selector-gradient problem.

Inference-time sparse selection. A parallel line keeps a dense-trained model and sparsifies only at inference, selecting a query-dependent subset of the KV cache: Quest [16] (top- k pages), H2O [17] (heavy-hitter eviction), and RetrievalAttention [18] (attention-aware ANN over keys, which reports that plain LSH/ANN underperforms because attention queries are out-of-distribution). LSH-E [19] uses SimHash for KV eviction, and DASH-KV [20] replaces LSH with a *learned* asymmetric hash, reporting higher retrieval recall than parameter-free LSH. Our train-dense / infer-LSH study (Section 5.6) sits in this regime, at block rather than token granularity and at small scale.

Selector training. NSA trains the selector by distilling dense per-block attention mass. MoBA uses a gating mechanism. Sparse Sinkhorn Attention [10] uses differentiable sorting via the Sinkhorn operator to learn latent permutations. We build on the MoE gate-coupling approach [6]: the selector score is injected as a log-sigmoid bias into the attention logits, creating a differentiable path. An auxiliary routing loss provides cold-start supervision.

Flash-style fusion. FlashAttention [1] tiles Q/K/V into blocks and computes online-softmax in SRAM, achieving exact attention with reduced memory I/O. FlashAttention-2 [2] improves parallelism and work partitioning. Our Triton kernel applies the same dataflow but iterates over *selected* key blocks via an index list rather than all blocks in sequence order.

Centroid routing. Routing Transformer [5] assigns tokens to k -means clusters. Reformer [9] uses locality-sensitive hashing. Our CentroidSSA routes blocks through learned centroids ($\mathcal{O}(N \cdot C)$) but uses InfoNCE contrastive training [14] instead of k -means, addressing the non-differentiability of hard cluster assignment.

Triton. Triton [13] is an intermediate language and compiler for tiled neural network computations, enabling custom GPU kernels at near-CUDA performance from Python. Our forward and backward kernels are implemented in Triton 3.7.

3 Method

3.1 Block-Sparse Attention with Gate Coupling

Given input $x \in \mathbb{R}^{B \times L \times D}$, we partition the sequence into $n_{\text{blk}} = L/B_s$ blocks of size B_s . Per layer:

1. **Projections.** $Q, K, V = \text{QKV}(x)$, reshaped to (B, H, L, d_h) .
2. **Block scores.** A small MLP selector computes block-level summaries:

$$s_q = \text{mean}(\text{sel_q}(x)_{\text{blk}}) \in \mathbb{R}^{B \times n_{\text{blk}} \times d_s} \quad (1)$$

$$s_k = \max(\text{mean}(\text{sel_k}(x)_{\text{blk}}), \max(\text{sel_k}(x)_{\text{blk}})) \quad (2)$$

$$S = \max(s_q s_{k, \text{mean}}^\top, s_q s_{k, \text{max}}^\top) \in \mathbb{R}^{B \times n_{\text{blk}} \times n_{\text{blk}}} \quad (3)$$

3. **Selection.** For each query block i , select the top- k key blocks by $S[i, \cdot]$, plus the own (diagonal) block. Causal masking excludes future blocks. Result: $\text{sel} \in \mathbb{Z}^{B \times n_{\text{blk}} \times k_{\text{tot}}}$.
4. **Gate coupling.** Add $\log\text{-sigmoid}(S[i, j])$ as a bias to the per-token attention logits for selected block j :

$$\text{logit}_{i,j} = \frac{Q_i K_j^\top}{\sqrt{d_h}} + \log\text{-sigmoid}(S_{i,j}) \quad (4)$$

This makes the output depend differentiably on S , so sel_q and sel_k receive task gradients.

5. **Attend.** Exact scaled-dot-product attention over the selected blocks only ($k_{\text{tot}} \cdot B_s$ keys per query block).

3.2 Auxiliary Routing Loss

For synthetic tasks (MQAR) where needle positions are known, we supervise the selector with cross-entropy to the true needle block:

$$\mathcal{L}_{\text{route}} = \text{CE}(\text{softmax}(S_{\text{last}}), \lfloor \text{needle_pos}/B_s \rfloor) \quad (5)$$

For centroid routing (Section 5), we replace this with an InfoNCE contrastive loss. The total loss is $\mathcal{L} = \mathcal{L}_{\text{task}} + \lambda_r \mathcal{L}_{\text{route}}$, with λ_r annealed to zero over the second half of training.

3.3 Complexity

- **Attention:** $\mathcal{O}(N \cdot k \cdot B_s)$ -linear in N .
- **Selection (block-pair):** $\mathcal{O}(n_{\text{blk}}^2) = \mathcal{O}((N/B_s)^2)$, quadratic on the reduced axis, shrunk by B_s^2 .
- **Selection (centroid):** $\mathcal{O}(n_{\text{blk}} \cdot C)$, linear.
- **Memory (fused kernel):** $\mathcal{O}(N)$ -only base Q/K/V, no gathered tensor.

4 Fused Triton Kernel

4.1 Forward Kernel

Each Triton program handles one (B, H, q_{blk}) tile. It loads the $B_s \times d_h$ query tile, then iterates over the k_{tot} selected key blocks via the index list sel . For each selected block:

1. Load K, V tiles ($B_s \times d_h$) from the KV cache at the indexed block offset.
2. Compute $B_s \times B_s$ attention scores with scale.
3. Apply causal mask (triu) for the own block; mask future blocks entirely.
4. Add gate bias $\log\text{-sigmoid}(S_{i,j})$ if enabled.
5. Update online-softmax state (running max m , running sum ℓ , accumulator acc) via the standard FlashAttention recurrence [1].

The output is acc/ℓ . No gathered K/V tensor or score matrix is ever materialized, all intermediate state lives in SRAM.

4.2 Backward Kernel

The backward pass follows the FlashAttention-2 recompute pattern [2] with three iterations over the selected blocks:

1. **Pass 1:** Recompute forward to obtain final m and ℓ .
2. **Pass 2:** Compute $D = \sum_j \text{rowsum}(P_j \cdot dP_j)$, the total diagonal term for the softmax Jacobian across all selected blocks.
3. **Pass 3:** Compute gradients:

$$dV = P^\top dO \quad (6)$$

$$dP = dO V^\top \quad (7)$$

$$dS = P \odot (dP - D) \cdot \text{scale} \quad (8)$$

$$dQ += dS K, \quad dK += dS^\top Q \quad (9)$$

dK and dV are written via `atomic_add` (multiple query blocks may select the same key block); dQ is written directly.

4.3 Integration

A custom `torch.autograd.Function` (`BSAttnFunction`) routes forward and backward through the Triton kernels. The `BlockSparseSSA(use_triton=True)` flag switches between the PyTorch reference and the fused path. Non-contiguous tensors from the QKV projection are made contiguous before the kernel call.

4.4 Verification

- **Forward:** 9.77×10^{-4} max diff vs. PyTorch SDPA (fp16, $B=2, H=4, L=64, d_h=32, k=3$, causal).
- **Backward:** dQ : 4.0×10^{-3} , dK : 4.4×10^{-3} , dV : 3.7×10^{-3} vs. PyTorch autograd (fp32, $k=3$, causal).
- **Integration:** forward 4.5×10^{-4} , gradient 1.1×10^{-3} vs. PyTorch path (same model, toggle `use_triton`).
- **Gate bias:** 9.77×10^{-4} with gate enabled.

5 Linear Selection: Centroid and LSH

Block-pair selection scales as $\mathcal{O}(n_{\text{blk}}^2)$. To remove this quadratic term, we propose CentroidSSA: route blocks through C learned centroids, analogous to k -means routing in the Routing Transformer [5] but with differentiable training.

5.1 Mechanism

Each block is assigned to its nearest centroid ($\text{argmax of } s_k \cdot \text{cent}^\top$). Each query block routes to its top- c_{pq} centroids and attends to all blocks in those buckets (capacity c_{cap} per bucket). Selection cost: $\mathcal{O}(n_{\text{blk}} \cdot C)$, linear.

5.2 The Circular Loss Failure

The original routing-alignment loss was:

$$\text{consensus} = \text{argmax}(s_q + s_k), \quad \mathcal{L} = \text{CE}(s_q, \text{consensus}) + \text{CE}(s_k, \text{consensus}) \quad (10)$$

This is *circular*: at cold start, s_q and s_k are near-uniform, so argmax is arbitrary noise. The loss collapsed to $\ln C$ (uniform), and retrieval hit stayed at 0.19.

5.3 InfoNCE Fix

We replace the circular loss with an InfoNCE contrastive objective [14]:

$$\mathcal{L}_{\text{InfoNCE}} = -\log \frac{\exp(\text{sim}(q_{\text{last}}, k_{\text{needle}})/\tau)}{\sum_j \exp(\text{sim}(q_{\text{last}}, k_j)/\tau)} \quad (11)$$

where $\text{sim}(q, k) = \text{softmax}(q/\tau) \cdot \text{softmax}(k/\tau)^\top$ and $\tau = 0.5$. The positive is the needle key block; all other key blocks are negatives. This is fully differentiable and has no argmax dependency.

5.4 Capacity Sweep

Table 1: Centroid capacity sweep (8-pair MQAR, 10k steps, curriculum). C = centroids, cap = bucket capacity.

C	cap	hit	acc
16	4	0.71	0.16
16	8	1.00	0.15
32	4	0.68	0.16
32	8	1.00	0.16
64	4	0.68	0.17
64	8	1.00	0.17
Block-sparse (baseline)		0.99	0.77

Table 1 shows that **capacity is the dominant variable**: cap = 4 $\rightarrow$ hit $\approx$ 0.70; cap = 8 $\rightarrow$ hit = 1.0. Centroid count (16 $\rightarrow$ 64) has negligible effect. InfoNCE + cap=8 achieves selector parity with block-sparse (hit=1.0), but task accuracy is stuck at $\sim$ 0.16 vs. block-sparse’s 0.77.

5.5 The Precision/Recall Bottleneck

With cap=8, each query block selects $c_{pq} \times \text{cap} + 1 = 17$ blocks. The needle is always in the set (hit=1.0) but surrounded by 12 extra distractor blocks. We tested three remedies:

- **Within-bucket refinement** (top- k from candidates): hit drops to 0.72, acc unchanged (0.17).
- **Train-high / inference-low** (train cap=8, eval cap=4): hit drops to 0.73, acc drops to 0.15.
- **Main q/k refinement** (score candidates with attention q/k instead of selector): hit 0.72, acc 0.17.

None closed the gap. The bottleneck is *not* routing precision but *representation quality*: the coarse centroid routing path shapes the attention’s q/k/v projections differently than block-sparse’s direct pairwise scoring, limiting token-level extraction even when the needle block is present.

5.6 LSH Bucketing: Train Dense, Infer Linear

Centroid routing loses information in the *representation* (queries are compared against lossy cluster summaries). A second linear-time selector, LSH bucketing [9], instead prunes the *search*: block representations are hashed by fixed random projections into buckets (n_{rounds} rounds), a query block reads its own bucket, and the candidates are re-scored with the full $s_q \cdot s_k$ cosine and truncated to block-sparse’s budget, so a selected block is scored at full fidelity. Selection is $\mathcal{O}(n_{\text{blk}} \cdot n_{\text{rounds}} \cdot (n_{\text{buckets}} + \text{cap}))$, linear in N ; representations are L2-normalized so routing, gate, and hash all act on the same angular quantity.

Under supervised routing, LSH nearly matches block-sparse and far exceeds centroid: MQAR accuracy 0.76 vs. 0.82 (block-sparse) vs. 0.17 (centroid) at $n_{\text{blk}}=16$, and 0.92 vs. 0.96 at $n_{\text{blk}}=64$, with recall plateauing at hit = 0.84 (an irreducible hash-collision miss). Linear selection *can* therefore preserve accuracy when the representation stays lossless, which the lossy centroid path cannot. This is the section’s main positive result.

The open part is training without supervision. The routing loss above uses the known needle block, available only on synthetic MQAR. We tested the unsupervised regime two ways. On MQAR under gate coupling alone, LSH is fragile: recall is highly sensitive to the cosine gate temperature (hit 0.09, 0.34, 0.02 at scale 5.66, 30, 50) and never reaches block-sparse’s gate-only level (acc 0.96, hit 0.78). On a real language model (TinyStories, 45M, gate coupling only) the result is sharper (Table 2): gate-only LSH reaches the *same* perplexity as own-block-only attention (no routing), and its perplexity does not improve with context length, whereas block-sparse beats the local baseline by 3.7 perplexity and improves with length. LSH attends to five blocks but adds nothing over attending to one: its content routing is inert.

Table 2: Gate-only TinyStories (45M, $B_s=16$, $L=512$, 3000 steps). Local-only is top- $k=0$ (own block, no routing). LSH matches the no-routing baseline; only block-sparse routing helps and improves with length.

	val PPL	PPL@128	PPL@256	PPL@512
Local-only (no routing)	21.43	20.7	20.6	21.7
LSH (linear routing)	21.49	21.0	21.0	21.7
Block-sparse routing	17.69	20.1	18.6	17.8

Across these runs one mechanism recurs: hard hashing gives the selector no gradient on blocks it does not already collide with, so when the selector is trained *through* the hash it cannot learn to route without supervision, while block-sparse’s all-pairs top- k scores every block on every step and learns to route in every regime.

This points to what the inference-cost-only objective permits: train through the dense all-pairs top- k selector (full gradient, identical to block-sparse) and swap in LSH bucketing only at inference, on the same normalized representations, so the hash never enters the gradient path.

For *retrieval* this works and scales. On MQAR, LSH inference recovers the dense-trained selector at a fixed inference budget, and recall holds (even improves) as context grows: at 8 rounds and 64 buckets (constant, so inference is $\mathcal{O}(n_{\text{blk}})$), $n_{\text{blk}}=128/256/512$ give hit 0.978/0.992/0.994 (Table 3). The single needle is reliably hashed into the query’s bucket.

For *language modeling* it does not. At $n_{\text{blk}}=128$ (block 4, $L=512$, gate only) and a matched attention budget (own block plus top-4), a dense selector that scores all blocks reaches 7.9 perplexity, but LSH selecting its top-4 from ~ 16 bucket candidates reaches 113 to 1115 (worse at fewer rounds). The dense selector at the same budget succeeds, so this is not a coverage floor: LSH’s bucketing fails to recall the handful of relevant blocks the dense selector finds. A single needle is easy to hash to; the blocks a language model needs per query are not, at least with this hash. The apparent unsupervised LM recovery held only at near-dense coverage (32 rounds for 32 blocks) and does not survive at a sparse budget.

So train-dense / infer-LSH yields a genuinely linear, accuracy-preserving selector for needle-style retrieval, but not yet for general language-model attention. The open problem is a hash whose recall of content-relevant blocks is high enough for the LM at a sparse budget; a bit-based SimHash [19] or a learned hash [20] (both used for token-level KV retrieval, the latter reporting higher recall than plain LSH) are the natural next steps for the block-level selector.

Table 3: Train dense, infer LSH at a *fixed* inference budget (8 rounds, 64 buckets, cap 8; $\mathcal{O}(n_{\text{blk}})$). Recall holds and improves as context grows. MQAR, supervised; oracle is the dense-trained selector.

n_{blk}	LSH acc	LSH hit	oracle acc
128	0.95	0.978	0.99
256	0.98	0.992	0.99
512	0.99	0.994	0.99

6 Language Model Experiments

6.1 Setup

- **Dataset:** TinyStoriesV2-GPT4 [12], 5M tokens, GPT-2 BPE tokenizer (vocab 50,257).
- **Model:** 6 layers, $d = 512$, $h = 8$, weight-tied embeddings. 44.9M (dense) / 45.1M (sparse) parameters.
- **Training:** 5000 steps, batch 16, $L = 512$, AdamW ($\beta_1=0.9, \beta_2=0.95, wd=0.1$), cosine LR 3×10^{-4} with 200-step warmup. bf16 mixed precision.
- **Sparse config:** $B_s = 64$, top- $k = 4$, gate coupling, Triton kernel.
- **Hardware:** NVIDIA GB10 Grace-Blackwell, 128 GB unified memory.

6.2 Results

Table 4: TinyStories language model: dense vs. block-sparse Triton.

Model	Params	Val PPL	PPL@128	PPL@256	PPL@512	Time
Dense	44.9M	14.41	15.41	14.54	14.09	736 s
Sparse-Triton	45.1M	15.00	15.31	14.16	14.51	908 s

The sparse-Triton model achieves val perplexity 15.0 versus 14.4 for dense, a 4% gap. Notably, at 256-token context, sparse *beats* dense (14.16 vs. 14.54), suggesting the learned selector provides a useful inductive bias at moderate lengths. The training time overhead (908s vs. 736s, +23%) is from the selector computation and the Triton kernel’s contiguity copies.

6.3 NIAH (Needle-in-a-Haystack)

Both models score 0% on NIAH at $L = 512$ (50 samples). This is a model capacity limitation (45M params, 3000 to 5000 steps), not a sparse attention defect. NIAH retrieval requires 100M+ parameters and 10k+ training steps.

7 Benchmark: Latency and Memory

Table 5: Forward-pass benchmark: $D=512, H=8, B=1, B_s=128, k=8$, causal. GB10 GPU.

L	Dense (ms)	Sparse-PT (ms)	Triton (ms)	Dense (GB)	PT (GB)	Triton (GB)
4,096	0.25	11.64	1.57	0.02	0.70	0.05
16,384	3.05	44.08	6.80	0.10	2.70	0.10
65,536	48.20	178.59	28.32	0.30	10.70	0.30
131,072	195.38	365.72	56.79	0.58	21.38	0.57
262,144	796.59	739.32	114.37	1.12	42.72	1.11

Memory. The Triton kernel uses 1.11 GB at 262k -parity with dense FlashAttention (1.12 GB) and $40\times$ less than the PyTorch gather reference (42.72 GB). The fused kernel never materializes the gathered K/V tensor; all intermediate state lives in SRAM.

Latency. The Triton kernel is $7.4\times$ faster than dense at 262k (114ms vs. 797ms). The crossover occurs at $\sim 32k$ tokens. At 65k, the kernel is already $1.8\times$ faster than dense.

PyTorch gather is the bottleneck. The same algorithm, same selection, but PyTorch materializes the gather (42.7 GB) and is $6.5\times$ slower than Triton at 262k. The fused kernel proves the sparse mechanism itself is strictly superior to dense at long contexts; the reference implementation was the obstacle.

7.1 Selection vs. the Attention Read: the $N^{4/3}$ Floor

The benchmark above measures the fused attention *read*, which is linear. The block-pair *selection* (scoring all n_{blk}^2 block pairs, Section 3) is not. We isolate the two stages and sweep L to 16M tokens across block sizes (Table 6).

Table 6: Selection (quadratic) vs. attention read (linear), isolated. $D=512$, $H=8$, $B_s=64$, $k=8$, causal, GB10 (130 GB). “OOM”: the $\mathcal{O}(n_{\text{blk}}^2)$ score matrix exceeds memory.

L	Selection (ms)	Read (ms)	Select mem (GB)
262,144	5.8	17.1	0.4
1,048,576	51.2	84.7	3.0
4,194,304	652.3	359.1	34.4
8,388,608	OOM	716.0	>130
16,777,216	OOM	1,447.6	>130

Scaling. Fitting across L , the attention read scales with exponent 0.96 to 1.02 (linear) at every block size, while selection scales with exponent 1.5 to 1.9 (quadratic, approaching 2 as N leaves the latency floor). Selection latency overtakes the read at $\sim 512\text{k}$ tokens for $B_s=32$ and $\sim 2\text{M}$ for $B_s=64$; for $B_s=128$ it stays below the read through 8M (ratio 0.17), crossing over only in the hundreds of millions.

The quadratic dies by OOM, not by latency. The $\mathcal{O}(n_{\text{blk}}^2)$ score matrix is a memory wall reached well before the latency crossover: it OOMs at 4M ($B_s=32$), 8M ($B_s=64$), and 16M ($B_s=128$), while the linear read runs to 16M (86 GB) on the same GPU. In practice, quadratic selection becomes impossible to hold, not merely slow.

Variable block size does not rescue linearity. A natural trick is to grow the block with context, $B_s \propto \sqrt{N}$, making selection $\mathcal{O}((N/\sqrt{N})^2) = \mathcal{O}(N)$, linear. We confirm this: along a $\sqrt{N}$ schedule ($B_s=32, 64, 128$ at $L=256\text{k}, 1\text{M}, 4\text{M}$) the measured selection exponent falls to 0.98. But the read processes kB_s keys per query, so it then grows as $NB_s \propto N^{1.5}$ (measured exponent rising to 2.3): the quadratic merely relocates from selection into the read. Writing the total as $N^2s/B_s^2 + NcB_s$ and minimizing over B_s gives $B_s \propto N^{1/3}$, at which *both* terms equal $\mathcal{O}(N^{4/3})$. This $N^{4/3}$ is a floor: **no block-size schedule makes single-level, all-pairs block selection linear end-to-end; the quadratic can be moved between stages but not removed.** Genuinely linear selection requires scoring $\mathcal{O}(C)$ candidates per query (clustering, hashing, or hierarchy), not all n_{blk} , the open problem of Section 5.

8 Discussion

8.1 What Works

Gate coupling + auxiliary routing loss reliably trains the block selector (hit 0.99 on MQAR, val PPL within 4% of dense on TinyStories [12]). The fused Triton kernel delivers the theoretical memory and latency advantages of block-sparse attention in practice, with no accuracy loss from numerical approximation.

8.2 What Doesn’t: Centroid Routing

Centroid routing achieves linear selection cost and perfect retrieval recall (hit=1.0) but cannot match block-sparse task accuracy (acc 0.16 vs. 0.77). The failure is not in routing (InfoNCE [14] fixes that) but

in *representation quality*: the coarse centroid path shapes the attention projections differently, limiting fine-grained token extraction. This is a fundamental limitation of content-routed linear-selection attention that, to our knowledge, has not been previously characterized.

8.3 Limitations

- Single-level block-pair selection is $\mathcal{O}(n_{\text{blk}}^2)$ with an $\mathcal{O}(N^{4/3})$ floor no block schedule can beat (Section 7.1); it exhausts memory before any latency crossover. Centroid routing is linear but accuracy-broken.
- NIAH requires a larger model; our 45M model is too small.
- The backward kernel uses `atomic_add` for dK/dV, which may serialize under high contention. A split-K or warp-level reduction could improve throughput.
- No autoregressive KV-cache support; the kernel is training/prefill only.

9 Conclusion

We demonstrated that block-sparse attention with a fused Triton kernel achieves memory parity with dense FlashAttention and $7.4\times$ latency speedup at long contexts, while maintaining within-4% perplexity on a real language model. The gate-coupled selector training recipe is simple, effective, and requires no separate pre-training stage.

We probed two linear-time selectors. Centroid routing reaches linear cost and perfect recall but degraded accuracy: the obstacle is representation quality under coarse routing, not routing recall (InfoNCE resolves the latter). LSH bucketing, which prunes the search rather than the representation, matches block-sparse accuracy under supervised routing and far exceeds centroid, showing a lossless-refine linear selector can preserve accuracy; but its unsupervised routing proved inert, matching a no-routing baseline on a 45M language model. Reliable unsupervised linear selection remains open.

Promising directions include scaling to the 100M+ regime where needle-in-a-haystack retrieval becomes measurable, extending the kernel to autoregressive KV-cache decoding, and designing a linear-selection mechanism that preserves the representation quality of pairwise scoring. More broadly, the fused kernel removes the implementation barrier that has masked block-sparse attention’s asymptotic advantage: at long contexts, content-dependent sparsity is now both faster and lighter than dense attention in practice, not merely in theory.

References

- [1] T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2205.14135.
- [2] T. Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. In *International Conference on Learning Representations (ICLR)*, 2024. arXiv:2307.08691.
- [3] J. Yuan, H. Gao, D. Dai, J. Luo, L. Zhao, Z. Zhang, Z. Xie, Y.-X. Wei, L. Wang, Z. Xiao, Y. Wang, C. Ruan, M. Zhang, W. Liang, and W. Zeng. Native sparse attention: Hardware-aligned and natively trainable sparse attention. *arXiv preprint* arXiv:2502.11089, 2025. ACL 2025 Best Paper.
- [4] E. Lu, Y. Duan, H. Hu, W. Wang, J. Zhu, H. Liu, K. Sun, and T. Wang. MoBA: Mixture of block attention for long-context LLMs. In *NeurIPS*, 2025. arXiv:2502.13189.
- [5] A. Roy, M. Saffar, A. Vaswani, and D. Grangier. Efficient content-based sparse attention with routing transforms. *Transactions of the Association for Computational Linguistics (TACL)*, 9:53-68, 2021.
- [6] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. V. Le, G. E. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *International Conference on Learning Representations (ICLR)*, 2017. arXiv:1701.06538.

- [7] I. Beltagy, M. E. Peters, and A. Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- [8] M. Zaheer, G. Guruganesh, A. Dubey, J. Ainslie, C. Alberti, S. Ontañón, et al. BigBird: Transformers for longer sequences. In *NeurIPS*, 2020. arXiv:2007.14062.
- [9] N. Kitaev, Ł. Kaiser, and A. Levskaya. Reformer: The efficient transformer. In *ICLR*, 2020. arXiv:2001.04451.
- [10] Y. Tay, D. Bahri, D. Metzler, D.-C. Juan, Z. Zhao, and C. Zheng. Sparse sinkhorn attention. In *ICML*, 2020. arXiv:2002.11296.
- [11] C. Olsson, N. Elhage, N. Nanda, N. Joseph, N. DasSarma, T. Henighan, B. Mann, et al. In-context learning and induction heads. *Transformer Circuits Thread*, Anthropic, 2022.
- [12] R. Eldan and Y. Li. TinyStories: How small can language models be and still speak coherent English? *arXiv preprint arXiv:2305.07759*, 2023.
- [13] P. Tillet, H.-T. Kung, and D. D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. In *Proc. 3rd ACM SIGPLAN Int. Workshop on Machine Learning and Programming Languages (MAPL)*, pages 10-19, 2019.
- [14] A. van den Oord, Y. Li, and O. Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [15] S. Arora, S. Gopi, N. Ingster, and P. Nakkiran. Simple, fast, and practically good: A simple LLM-based sparse retrieval model for MQAR. 2024.
- [16] J. Tang, Y. Zhao, K. Zhu, G. Xiao, B. Kasikci, and S. Han. Quest: Query-aware sparsity for efficient long-context LLM inference. In *ICML*, 2024. arXiv:2406.10774.
- [17] Z. Zhang, Y. Sheng, T. Zhou, T. Chen, L. Zheng, R. Cai, Z. Song, Y. Tian, C. Ré, C. Barrett, Z. Wang, and B. Chen. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *NeurIPS*, 2023. arXiv:2306.14048.
- [18] D. Liu, M. Chen, B. Lu, H. Jiang, Z. Han, Q. Zhang, Q. Chen, C. Zhang, B. Ding, K. Zhang, C. Chen, F. Yang, Y. Yang, and L. Qiu. RetrievalAttention: Accelerating long-context LLM inference via vector retrieval. *arXiv preprint arXiv:2409.10516*, 2024.
- [19] T. Rabbani, M. Liu, T. O’Halloran, A. Sankaralingam, M.-A. Hartley, and F. Huang. LSH-E: An adaptive locality-sensitive strategy for KV cache compression. In *NeurIPS*, 2024.
- [20] J. Guo, Z. Zhang, J. Xie, M. T. Iqbal, D. Han, L.-H. Lee, S.-H. Bae, J. Zou, Y. Yang, and C. Zhang. DASH-KV: Accelerating long-context LLM inference via asymmetric KV cache hashing. *arXiv preprint arXiv:2604.19351*, 2026.